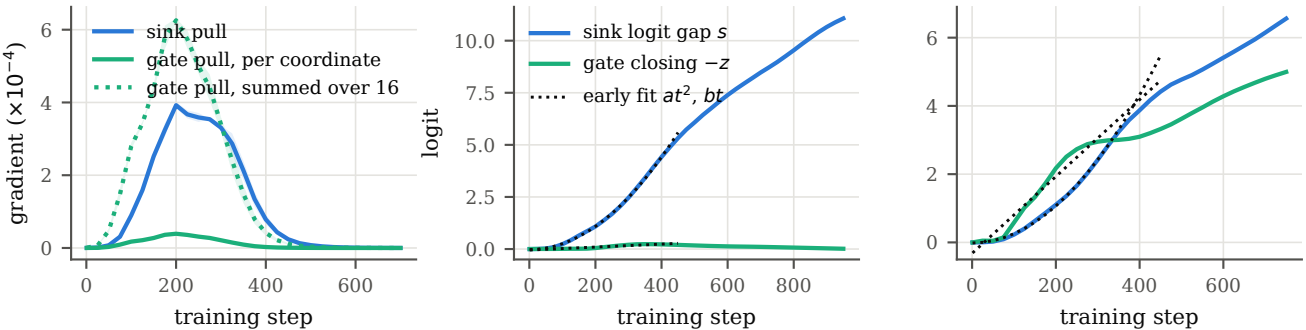
(a) Which way the loss pulls
gradient (×10⁻⁴)
sink pull
gate pull, per coordinate
gate pull, summed over 16
training step
(b) λ=1: the sink wins
logit
sink logit gap s
gate closing −z
early fit at², bt
training step
(c) λ=10: both routes share the work
training step